

ECE695 – ILGM

1. Kalman Filter with Neural Networks
2. Variational Autoencoders with Normalizing Flows

By: Sheikh Shams Azam

04/26/2021

Kalman Filter with Nueral Network

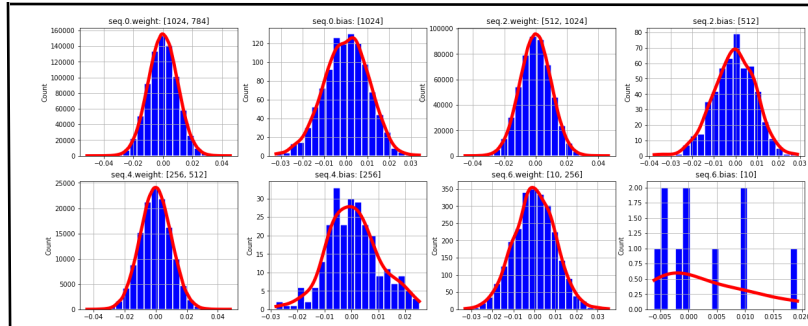
Motivation

- Current research in studying the optimization subspace in gradient descent.
- Came across this paper [1], One line summary:
"gradient descent tell us more than what we care about".
- Idea:
 - Neural Networks (NN) fail with confidence, i.e., false predictions with high confidence.
 - Solutions to this generally involves ensembling of multiple models trained independently.
 - In the optimization process of NNs, we discard the steps taken and only store the end weights.
 - Proposal: The path taken by the optimization can help us learn distribution of weights.
 - The parameter distribution can help us draw several models.
 - So basically, we can try to do a Bayesian prediction using NNs.

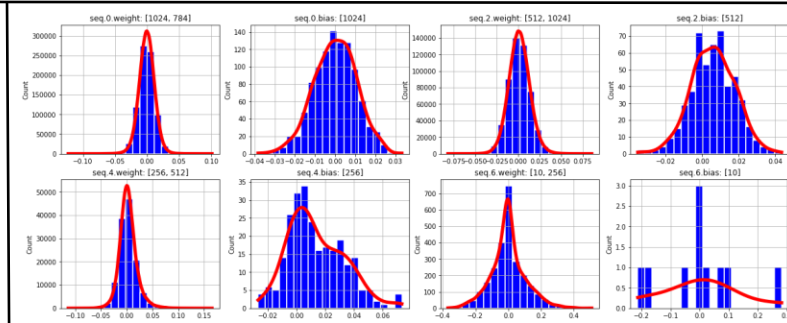
$$\underbrace{q(X_{n+1}|\theta^*)}_{\text{Frequentist}} \text{ vs } \underbrace{\int_{\theta} q(X_{n+1}|\theta)q(\theta|x)d\theta}_{\text{Bayesian}}$$

Neural Network Training


- Often NNs initialized using standard Gaussian distributions, $\theta_0 \sim \mathcal{N}(\mu_{\theta_0}, \sigma_{\theta_0}^2)$.
- Training NN with gradient descent is a linear update: $\theta_{k+1} = \theta_k - \eta \cdot \nabla f_{\theta_k}(X, Y)$.
- A Gaussian distributed RV undergoing a linear (affine) transformation gives another Gaussian.



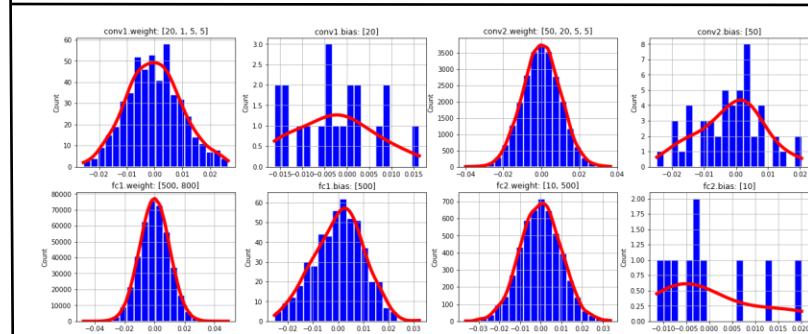
FCN: initialization



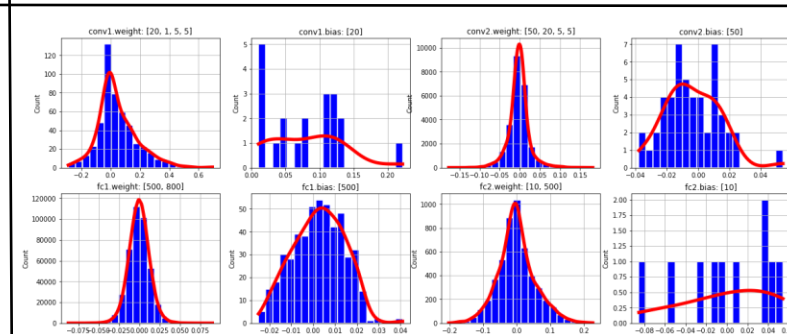
FCN: epoch 20



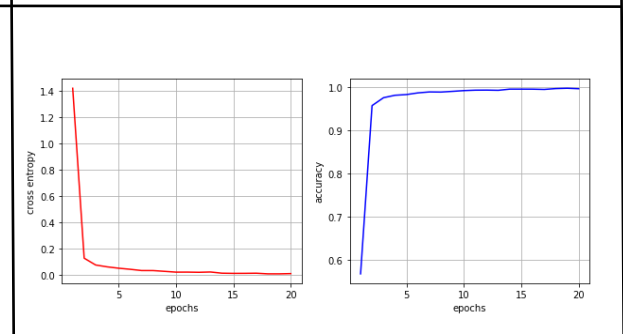
FCN: training



CNN: initialization



CNN: epoch 20



CNN: training

HMM of Weights: Kalman Filter


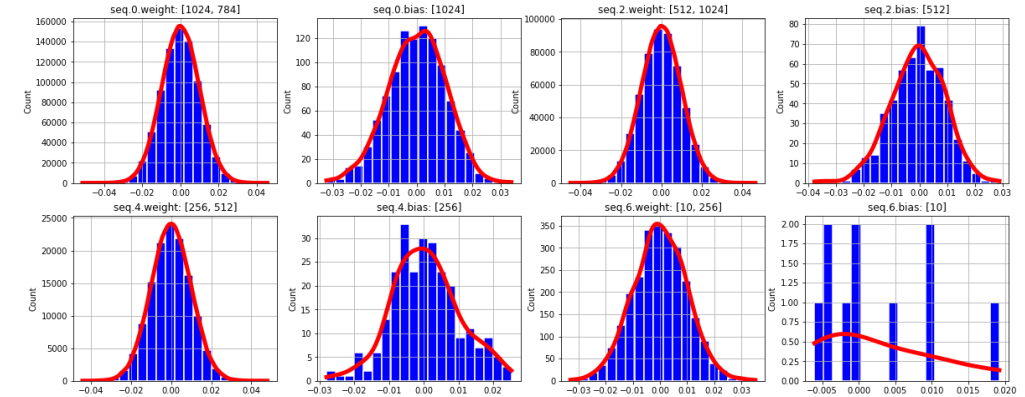
- Means of the weights follow a Markov process:

$$\mu_{k+1} = \mu_k - \mathbb{E}[\eta \cdot \nabla f_{\theta_k}(X, Y)]$$

- Similarly, variance of weights can be modelled:


$$\sigma_{k+1}^2 = \sigma_k^2 + \eta^2 \mathbb{E}[(\nabla f_{\theta_k})^2] - \eta^2 \mathbb{E}^2[\nabla f_{\theta_k}]$$

- Tracking the means becomes a direct application of Kalman filtering in this case.
- Since the emissions are Gaussian distributed we get a recursive solution for tracking the means and variance.
- Simplifying assumptions:
 - Weights in different layers are independent of each other
 - Weights within the same layer can be considered pairwise related in which case it is possible to track the covariance matrix.



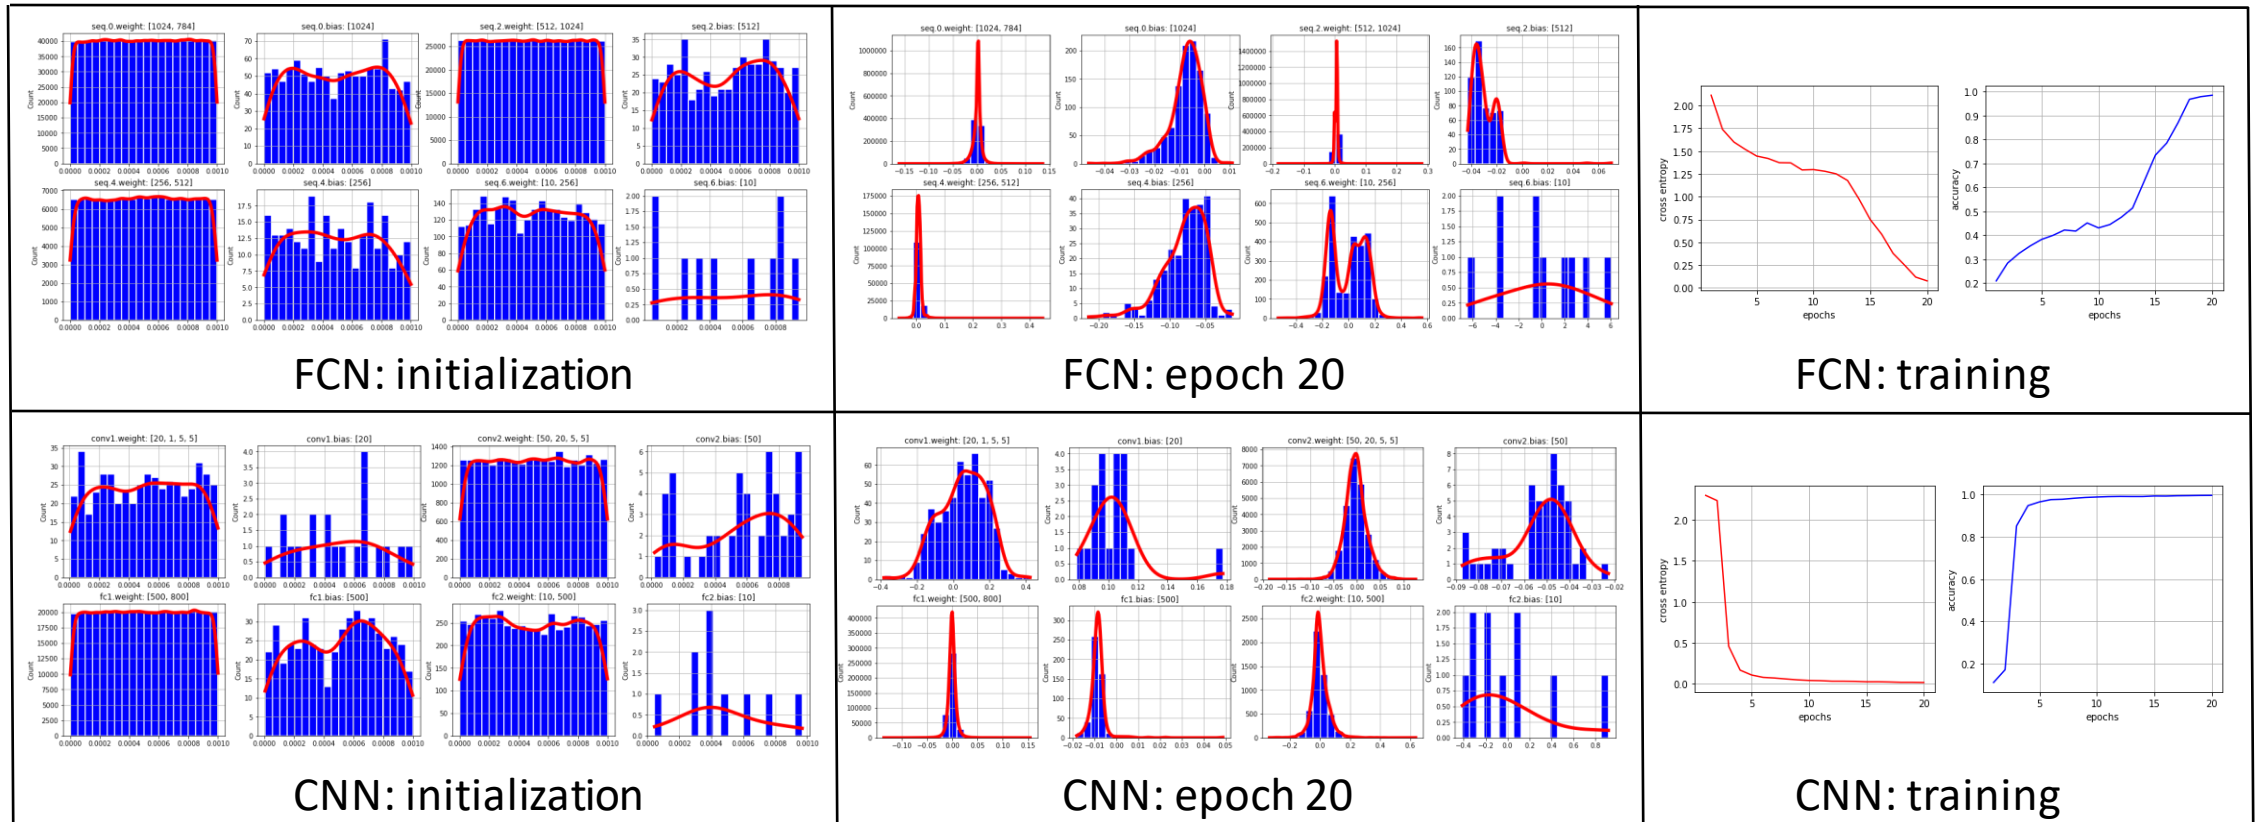
Some Results: Kalman Filter

- Trained model gives the frequentist prediction
- Models drawn from the probability distribution give better picture about the uncertainty in prediction.
- Classification on noisy data (bottom graph) gives a good picture of uncertainty in prediction.
- Can be used to detect out of distribution samples by checking the variance in prediction.
- There are other methods like deep-ensemble, MC Dropout methods in the same domain.
- Memory requirement of saving ensembles is higher than weight distribution tracking.



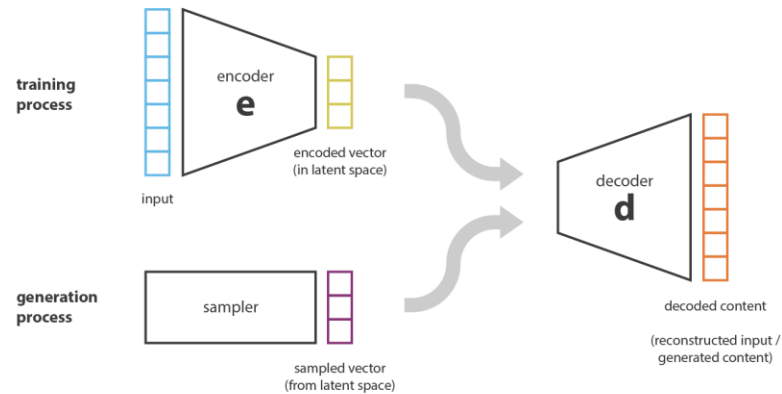
Non-Gaussian Weight Initialization

- NNs can also be initialized using non-Gaussian distributions.
- Simple case of uniform distribution. Final distributions are non-Gaussian like.
- Kalman-Filter could be replaced with Particle filter instead but computation overhead.



Normalizing Flow with Variational Autoencoder

Normalizing Flows with VAE




$$\underbrace{\mathbb{E}_{r(X|Y)}[-\log q(X|Y)] - D_{KL}[r(X|Y; \phi) || q(X; \theta)]}_{\text{Reconstruction loss}}$$

Reconstruction loss

VAE: The posterior is approximated with a shifted and scaled Gaussian (decouples sampling and X).

Normalizing flows was originally proposed to remove this simplification!!!



References:

- Implementation for this work (final update soon): <https://github.com/shams-sam/EE695-ILGM>.
- Kalman and Bayesian Filters in Python: <https://github.com/rlabbe/Kalman-and-Bayesian-Filters-in-Python> is a good resource for introduction to other filtering approaches.
- Franchi, Gianni, et al. "TRADI: Tracking deep neural network weight distributions." *European Conference on Computer Vision (ECCV) 2020*. (code for this paper is not released; vectorized implementation is available in my project: <https://github.com/shams-sam/EE695-ILGM>).
- Kingma, Diederik P., and Max Welling. "Auto-encoding variational bayes." *arXiv preprint arXiv:1312.6114* (2013).
- Kobyzev, Ivan, Simon Prince, and Marcus Brubaker. "Normalizing flows: An introduction and review of current methods." *IEEE Transactions on Pattern Analysis and Machine Intelligence* (2020).